

# Bayesian ancestral state reconstruction

*Simon Joly*

*BIO 6008 - Fall 2015*

## Contents

<b>Bayesian ancestral state reconstructions</b>	<b>1</b>
<b>Multistate reconstruction</b>	<b>2</b>
Running the analysis . . . . .	2
Bayesian analysis diagnostics . . . . .	4
<b>Corelated evolution between binary traits in BayesTraits</b>	<b>12</b>
Independent model . . . . .	12
Dependent model . . . . .	13
Running the analysis . . . . .	13
Bayes Factors . . . . .	19
<b>Assignment</b>	<b>20</b>
<b>References</b>	<b>20</b>

## Bayesian ancestral state reconstructions

In the previous lecture, we saw how to reconstruct ancestral state reconstruction using maximum likelihood or stochastic mapping, the latter of which uses monte carlo simulations. In this lecture, we will see how to reconstruct ancestral states using a full Bayesian approach (Pagel et al. 2004).

For this, we will use the program BayesTraits, written by Andrew Meade and Mark Pagel. Conveniently, Randi Griffin has written wrapper functions that allow to call BayesTrait from R. The BayesTrait wrapper function can be downloaded from her [website](#). However, I have included modified scripts with this lecture of the BayesTraits Wrapper that include several important improvements. For instance, I modified the scripts to make the wrapper multi-platform (the original was only for OSX) and to use the second version of BayesTraits. The BayesTrait Wrapper can be found in the folder BTW.

### Prepare seed plant data

Throughout this tutorial, we will use the seed plant phylogeny and trait data from Paquette et al. (2015). Let's load it and prepare it.

```
# Load ape
require(ape)
# Import datasets
seedplantstree <- read.nexus("./data/seedplants.tre")
```

```

seedplantsdata <- read.csv2("./data/seedplants.csv")
# Remove species for which we don't have complete data
seedplantsdata <- na.omit(seedplantsdata)
# Remove species in the tree that are not in the data matrix
species.to.exclude <- seedplantstree$tip.label[!(seedplantstree$tip.label %in%
                                                seedplantsdata$Code)]
seedplantstree <- drop.tip(seedplantstree,species.to.exclude)
rm(species.to.exclude)
# Name the rows of the data.frame with the species codes used as tree labels
rownames(seedplantsdata) <- seedplantsdata$Code
seedplantsdata <- seedplantsdata[,-1]
# Order the data in the same order as the tip.label of the tree. In the present
# example, this was already the case.
seedplantsdata <- seedplantsdata[seedplantstree$tip.label,]
# Create a factor for a categorical variable
height <- factor(seedplantsdata$height)
names(height) <- rownames(seedplantsdata)
# Create a vector for a continuous character
maxH <- seedplantsdata$maxH
names(maxH) <- rownames(seedplantsdata)

```

## Multistate reconstruction

We will use the MultistateMCMC R function to estimate ancestral states using a Bayesian approach in BayesTraits (Pagel et al. 2004). The first thing to do is to load all the functions from BayesTrait Wrapper. Copy the BTW folder into your working folder and then enter the following code:

```

for (n in 1:length(list.files('./BTW/R'))){
  source(paste("./BTW/R/", list.files('./BTW/R')[n], sep=""))
}

```

You will then have to copy the “BayesTraits” programs into your working directory. You can only copy the program that corresponds to your operating system.

## Running the analysis

BayesTraits has several functions that can be modified. The most important are implemented in the BTW functions, but maybe not all of them. For a complete description of the functions available in these function, you should have a look at the BTW manual in the folder `./BTW/help/BTWman.pdf`.

When running a Bayesian analysis of BayesTraits from R, the following parameters are important for Bayesian MCMC analyses:

Parameter	Description
it	integer specifying number of MCMC iterations. Default is 10000.
bi	integer specifying number of iterations to discard as burn-in. Default is 1000.
sa	integer specifying number of iterations to skip between samples. Default is 100.
rd	positive number specifying the rate deviation parameter. Default is 2.

When running the a Multistate analysis, these options are also important:

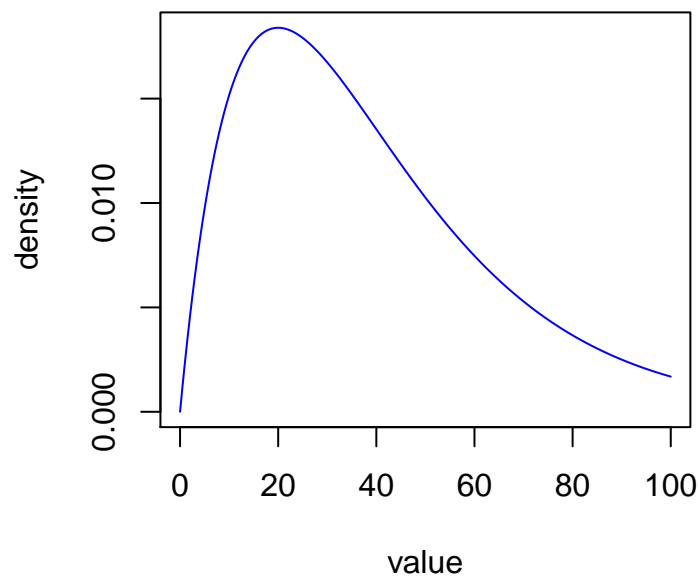
Parameter	Description
res	character or vector indicating restrictions to place on rates. If a vector is given, each element indicates an index
resall	character indicating a rate or a non-negative number to restrict all rates to.
mrca	character or vector indicating nodes to reconstruct using the most recent common ancestor approach. If a vector is given, each element indicates a node
fo	character or vector indicating nodes to fix at particular states. If a vector is given, each element is a character state
et	character or vector listing taxa to exclude.

And finally, the MultistateMCMC analysis have these additional parameters:

Parameter	Description
pr	character or vector describing prior distributions for model parameters. If a vector is given, each element is a character string
pa	character string specifying the prior distribution for all parameters by listing first the name of the parameter, then the prior distribution
rj	toggles reversible jump model if non-empty. A character string specifying the prior distribution and it's parameters
rjhp	toggles reversible jump model with a hyper-prior if non-empty. A character string specifying the prior distribution and hyper-prior
hp	character or vector describing prior distributions and hyper-priors for model parameters. If a vector is given, each element is a character string
hpall	character string specifying the prior distribution and hyper-prior for all parameters by listing first the name of the parameter, then the prior distribution

We will now run three independent MCMC runs of BayesTraits. This is important to make sure that the analyses have converged on the same estimates. We will use the same substitution model as last week, that is with three different rates. We will run an analysis of 100000 generations (`it`), sampling the chain every 100 generations (`sa=100`), discarding the first 1000 as burnin (`bi`). Finally, a gamma prior with mean 2 and shape 20 will be given to all parameters (`pa="gamma 2 20"`). This distribution looks like the following:

```
x <- seq(0, 100, length=200)
hx <- dgamma(x, shape=2, scale=20)
plot(x, hx, type="l", ylab="density", xlab="value", col="blue")
```



Now, let's run BayesTraits.

```

height.dat<-data.frame(code=rownames(seedplantsdata),height=as.numeric(seedplantsdata$height))
# Model: rate constraints
constraints <- c("q21 q13", "q23 q32", "q12 q31")
# Run three independent analyses (MCMC chains)
multistate.MCMC.res1 <- MultistateMCMC(seedplantstree, height.dat, res=constraints,
  it = 100000, bi = 1000, sa = 100, rd = 2, pa = "gamma 2 20", silent = TRUE)
multistate.MCMC.res2 <- MultistateMCMC(seedplantstree, height.dat, res=constraints,
  it = 100000, bi = 1000, sa = 100, rd = 2, pa = "gamma 2 20", silent = TRUE)
multistate.MCMC.res3 <- MultistateMCMC(seedplantstree, height.dat, res=constraints,
  it = 100000, bi = 1000, sa = 100, rd = 2, pa = "gamma 2 20", silent = TRUE)

```

## Bayesian analysis diagnostics

For diagnostic of Bayesian MCMC analyses, the package `coda` is very useful to look for chain convergence and calculate statistics. To be able to estimate convergence statistics, it is important to run at least 2 independent chains. This should be standard anyway to ensure convergence and thus that the results are reliable. First, let's convert the output of `BayesTraits` into `coda` format.

```

require(lattice)
require(coda)
# Read the BayesTrait results in coda format
res1 <- mcmc(multistate.MCMC.res1$Results[,c(-1,-4)],
  start=min(multistate.MCMC.res1$Results$Iteration),
  end=max(multistate.MCMC.res1$Results$Iteration),thin=100)
res2 <- mcmc(multistate.MCMC.res2$Results[,c(-1,-4)],
  start=min(multistate.MCMC.res2$Results$Iteration),
  end=max(multistate.MCMC.res2$Results$Iteration),thin=100)
res3 <- mcmc(multistate.MCMC.res3$Results[,c(-1,-4)],
  start=min(multistate.MCMC.res3$Results$Iteration),
  end=max(multistate.MCMC.res3$Results$Iteration),thin=100)
# Combine the three chains
res <- mcmc.list(res1,res2,res3)

```

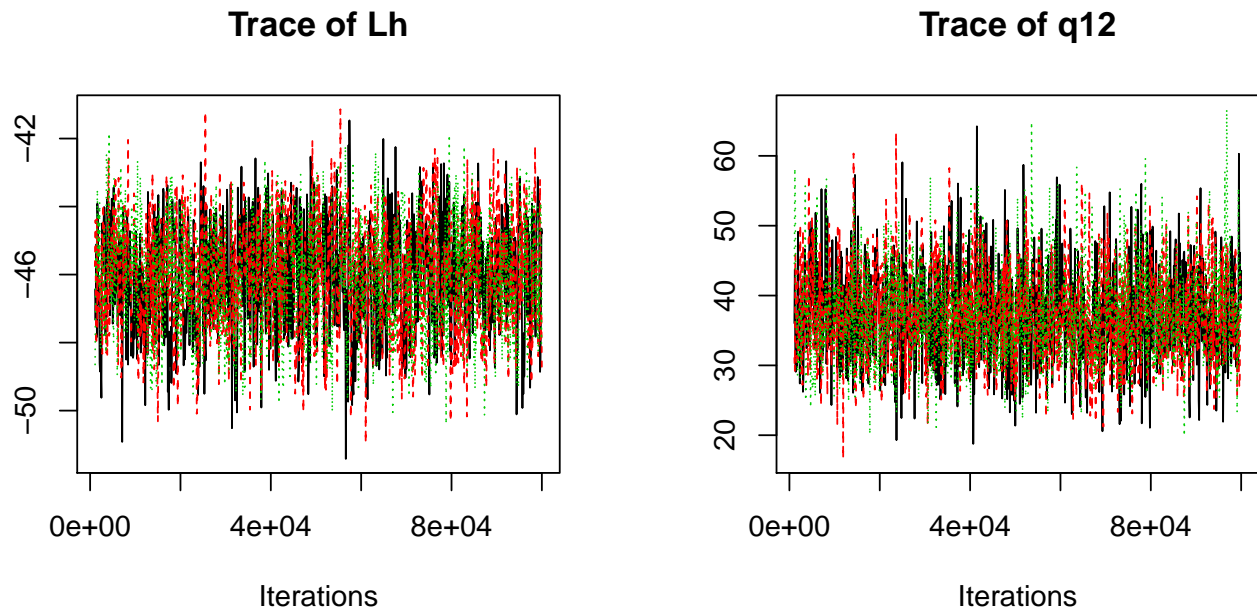
## Trace plots

Now, we can have a look at the results. Let's start by looking at the values of two parameters along the MCMC chain.

```

# Look at the trace plots for some characters
op <- par(mfrow=c(1,2))
traceplot(res[,c(1,3)])

```



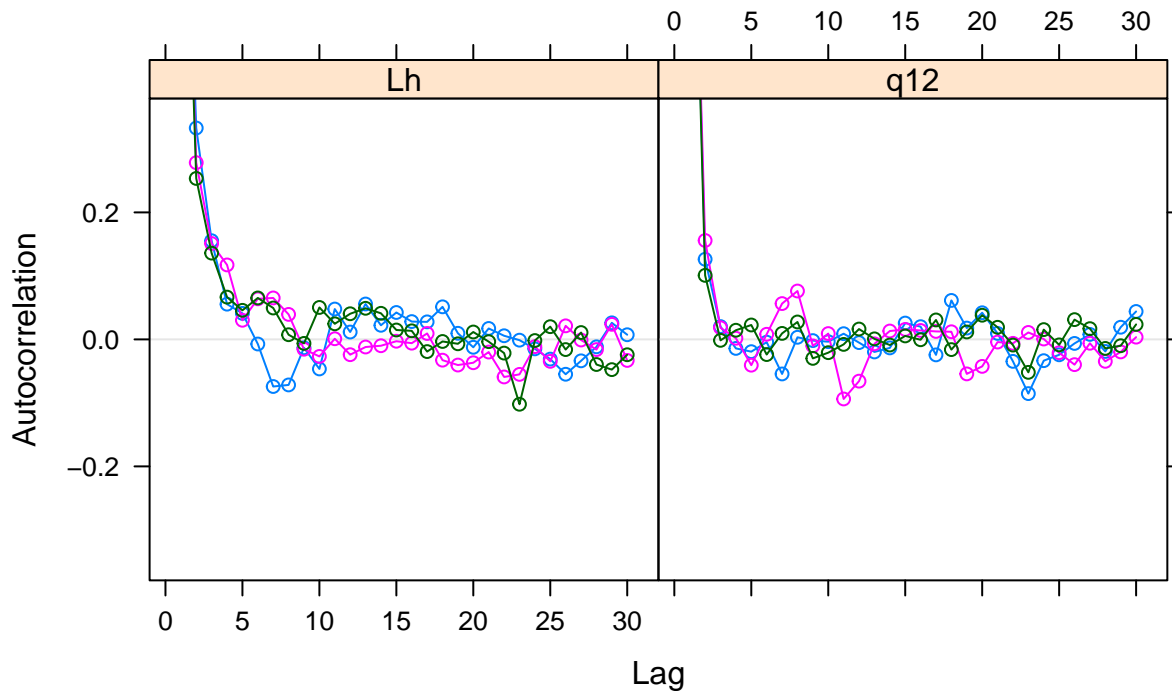
```
par(op)
```

The different colors on the plot represent the different chains. You can see that the values go up-and-down a lot, which is a sign that the chain is mixing well. The opposite would give a lot of correlations between successive samples and would give poor estimates of the parameters.

### Autocorrelation plots

You can see how the correlation drops between successive samples by using the function `acfplot`.

```
acfplot(res[,c(1,3)])
```



You can see that when samples are approximately 5 samples apart, they are not much correlated.

### Convergence diagnostics

Let's now look at some convergence diagnostics. The effective size of the parameter represents the estimated number of independent samples that are used to estimate the parameter's mean. Because parameter values are sampled from a chain, values sampled consecutively along the chain are generally correlated. The effective size is the estimated number of independent samples remaining once that autocorrelation is removed (this is inferred, of course). You generally want to have at least 200 of effective size to believe in your results (the more the better).

```
# Get effective sizes (should be > 200)
effectiveSize(res)
```

```
##           Lh Harmonic.Mean           q12           q13           q21
##    1412.1418       75.8133    2332.0987    1370.3311    1370.3311
##           q23           q31           q32    Root.P.1.    Root.P.2.
##    1160.5912    2332.0987    1160.5912    2355.2098    2582.5137
##    Root.P.3.
##    2524.7686
```

The Gelman and Rubin's Potential Scale Reduction Factor (PSRF) is based on a comparison of within-chain vs. between-chain variance. If the chains have converged, then the potential scale reduction factor should be 1. If the values are above 1.05, this means you should run the chains longer.

```
# Gelman and Rubin's convergence diagnostic
gelman.diag(res, autoburnin=FALSE, multivariate=FALSE)
```

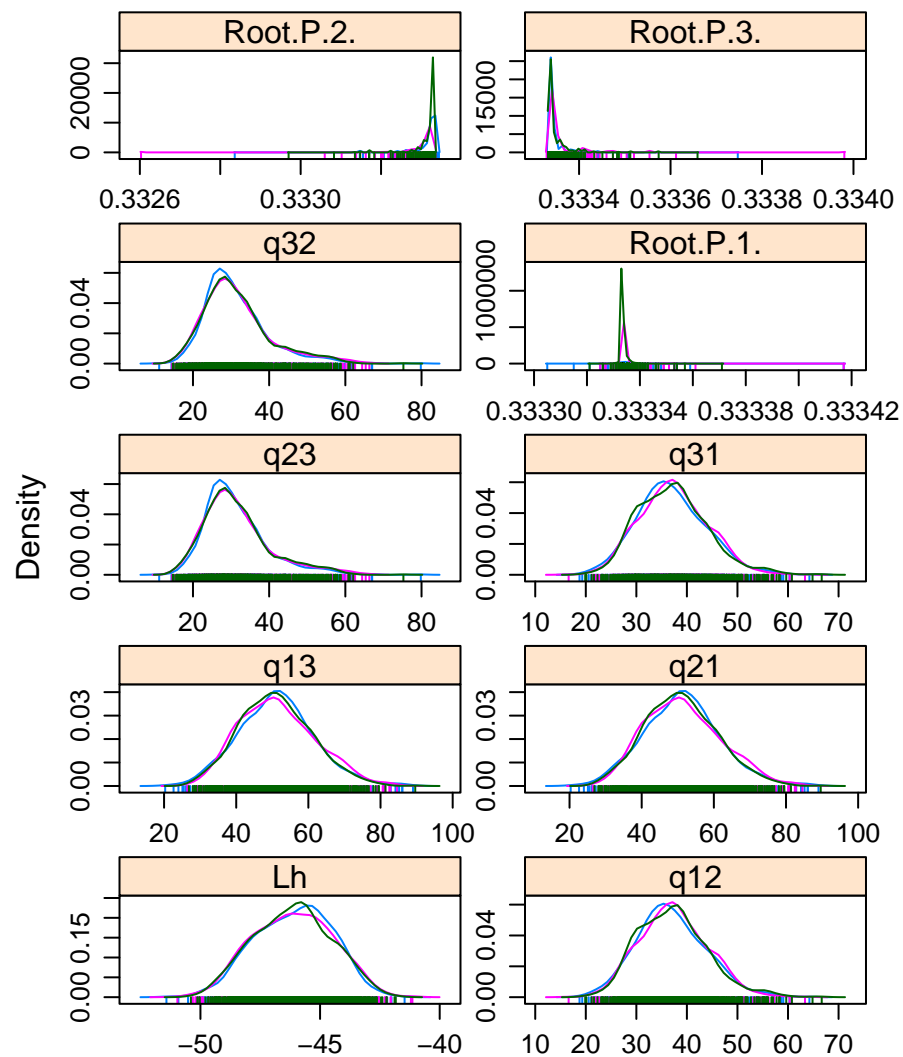
```
## Potential scale reduction factors:
```

```
##
##          Point est. Upper C.I.
## Lh          1.00      1.00
## Harmonic.Mean 1.41      2.23
## q12          1.00      1.00
## q13          1.00      1.00
## q21          1.00      1.00
## q23          1.00      1.00
## q31          1.00      1.00
## q32          1.00      1.00
## Root.P.1.    1.02      1.02
## Root.P.2.    1.03      1.03
## Root.P.3.    1.04      1.04
```

## Density plots

Now, let's look at the density plots for the parameters.

```
# Density Plots
densityplot(res[, -2])
```



```
# Parameter summary
summary(res)
```

```
##
## Iterations = 1100:1e+05
## Thinning interval = 100
## Number of chains = 3
## Sample size per chain = 990
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##              Mean          SD Naive SE Time-series SE
## Lh            -46.1236 1.627e+00 2.986e-02 4.351e-02
## Harmonic.Mean -47.3110 1.939e-01 3.558e-03 4.518e-02
## q12           37.1792 6.822e+00 1.252e-01 1.413e-01
## q13           51.2271 1.033e+01 1.896e-01 2.803e-01
## q21           51.2271 1.033e+01 1.896e-01 2.803e-01
## q23           31.4778 8.719e+00 1.600e-01 2.594e-01
## q31           37.1792 6.822e+00 1.252e-01 1.413e-01
## q32           31.4778 8.719e+00 1.600e-01 2.594e-01
## Root.P.1.      0.3333 3.192e-06 5.858e-08 6.550e-08
## Root.P.2.      0.3333 2.517e-05 4.618e-07 4.981e-07
## Root.P.3.      0.3333 2.274e-05 4.172e-07 4.600e-07
##
## 2. Quantiles for each variable:
##
##              2.5%      25%      50%      75%      97.5%
## Lh            -49.2245 -47.3376 -46.0729 -44.9554 -43.1323
## Harmonic.Mean -47.5415 -47.4162 -47.3409 -47.2902 -46.7293
## q12           25.3892 32.4744 36.8237 41.5310 51.3589
## q13           32.7544 43.9004 50.7917 57.8294 72.8183
## q21           32.7544 43.9004 50.7917 57.8294 72.8183
## q23           18.7282 25.6041 29.8007 35.4955 54.3993
## q31           25.3892 32.4744 36.8237 41.5310 51.3589
## q32           18.7282 25.6041 29.8007 35.4955 54.3993
## Root.P.1.      0.3333 0.3333 0.3333 0.3333 0.3333
## Root.P.2.      0.3333 0.3333 0.3333 0.3333 0.3333
## Root.P.3.      0.3333 0.3333 0.3333 0.3333 0.3334
```

```
# Highest Posterior Density intervals
HPDinterval(res)
```

```
## [[1]]
##              lower      upper
## Lh            -48.996878 -43.165809
## Harmonic.Mean -47.608705 -47.338061
## q12           23.021329 49.639770
## q13           31.690352 72.956165
## q21           31.690352 72.956165
## q23           15.911592 48.051960
## q31           23.021329 49.639770
## q32           15.911592 48.051960
```



```
## Root.P.1.      0.333332  0.333337
## Root.P.2.      0.333311  0.333334
## Root.P.3.      0.333332  0.333353
## attr("Probability")
## [1] 0.9494949
##
## [[2]]
##               lower      upper
## Lh             -49.254204 -43.200092
## Harmonic.Mean -47.430681 -46.921486
## q12            25.048473  49.594601
## q13            33.471411  72.818220
## q21            33.471411  72.818220
## q23            15.597003  50.930377
## q31            25.048473  49.594601
## q32            15.597003  50.930377
## Root.P.1.      0.333333  0.333337
## Root.P.2.      0.333306  0.333334
## Root.P.3.      0.333333  0.333359
## attr("Probability")
## [1] 0.9494949
##
## [[3]]
##               lower      upper
## Lh             -49.344520 -43.283237
## Harmonic.Mean -47.368667 -46.691168
## q12            25.232160  50.992961
## q13            32.689210  70.929942
## q21            32.689210  70.929942
## q23            16.265043  50.467956
## q31            25.232160  50.992961
## q32            16.265043  50.467956
## Root.P.1.      0.333331  0.333337
## Root.P.2.      0.333305  0.333334
## Root.P.3.      0.333333  0.333359
## attr("Probability")
## [1] 0.9494949
```

The density plots show that the runs have converged on very similar posterior distributions, which confirms the convergence diagnostic stats. The summary gives the quantiles and the median value.

Interestingly, whereas the transition rate parameters are in the same order as with likelihood inference, the variation is much smaller... Actually, this is a consequence of the prior used for the rate variation. If we take a flat prior instead between 0 and 1000, here is what we would get:

```
multistate.MCMC.res1 <- MultistateMCMC(seedplantstree, height.dat, res=constraints,
  it = 100000, bi = 1000, sa = 100, rd = 2, pa = "uniform 0 1000", silent = TRUE)
multistate.MCMC.res2 <- MultistateMCMC(seedplantstree, height.dat, res=constraints,
  it = 100000, bi = 1000, sa = 100, rd = 2, pa = "uniform 0 1000", silent = TRUE)
multistate.MCMC.res3 <- MultistateMCMC(seedplantstree, height.dat, res=constraints,
  it = 100000, bi = 1000, sa = 100, rd = 2, pa = "uniform 0 1000", silent = TRUE)
# Read the BayesTrait results in coda format
res1 <- mcmc(multistate.MCMC.res1$Results[,c(-1,-4)],
  start=min(multistate.MCMC.res1$Results$Iteration),
```

```

        end=max(multistate.MCMC.res1$Results$Iteration),thin=100)
res2 <- mcmc(multistate.MCMC.res2$Results[,c(-1,-4)],
            start=min(multistate.MCMC.res2$Results$Iteration),
            end=max(multistate.MCMC.res2$Results$Iteration),thin=100)
res3 <- mcmc(multistate.MCMC.res3$Results[,c(-1,-4)],
            start=min(multistate.MCMC.res3$Results$Iteration),
            end=max(multistate.MCMC.res3$Results$Iteration),thin=100)
# Combine the three chains
res <- mcmc.list(res1,res2,res3)
# Get effective sizes (should be > 200)
effectiveSize(res)

```

```

##          Lh Harmonic.Mean          q12          q13          q21
##    37.229224    10.509809    25.170421    8.534771    8.534771
##          q23          q31          q32    Root.P.1.    Root.P.2.
##    119.025292    25.170421    119.025292    1739.001367    1761.358625
##    Root.P.3.
##    1727.878694

```

```

# Gelman and Rubin's convergence diagnostic
gelman.diag(res,autoburnin=FALSE,multivariate=FALSE)

```

```

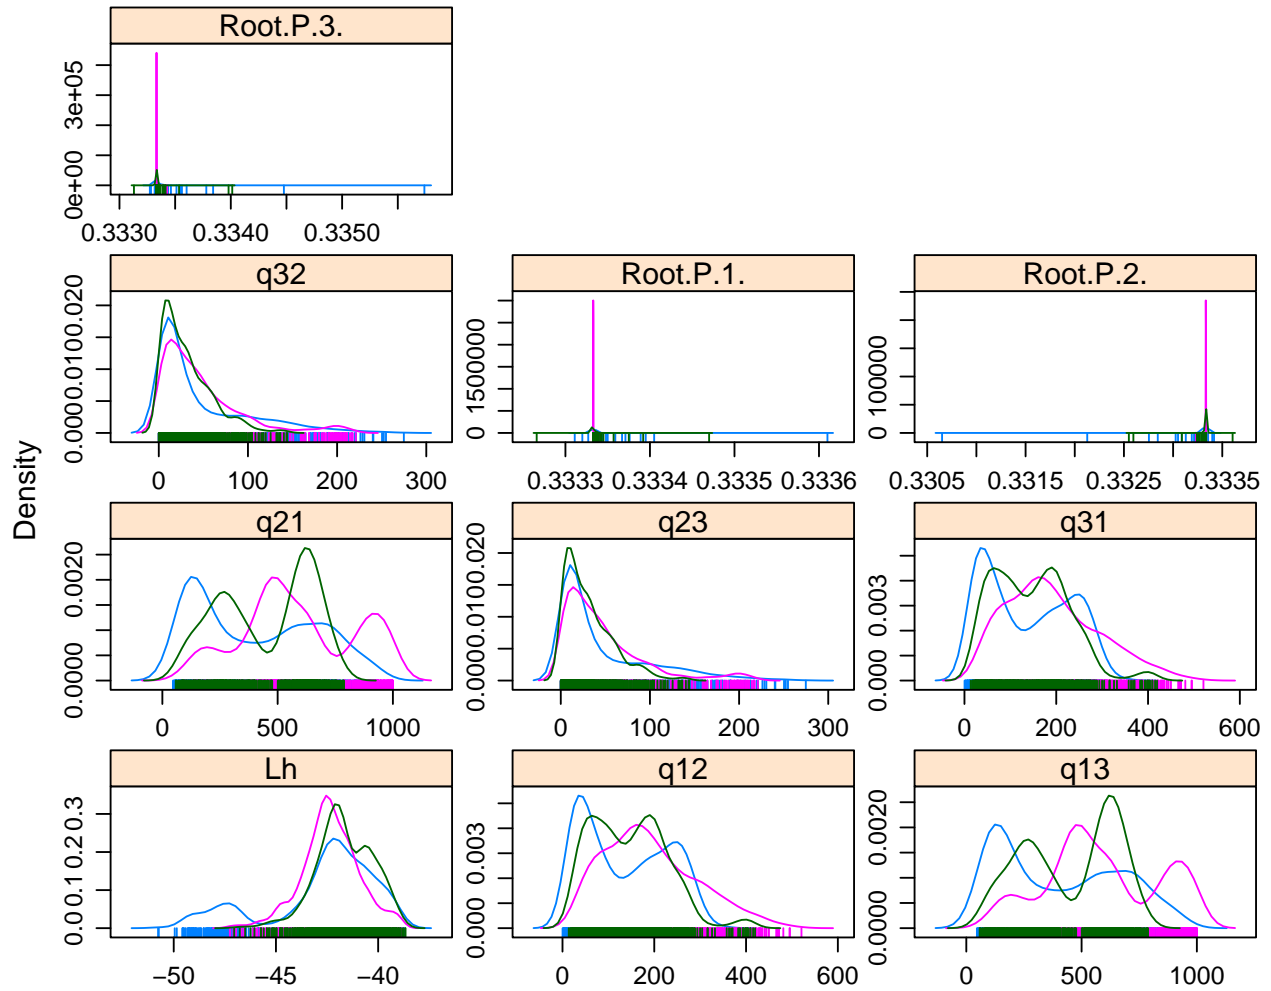
## Potential scale reduction factors:
##
##          Point est. Upper C.I.
## Lh          1.15      1.39
## Harmonic.Mean 6.36    12.82
## q12          1.09      1.26
## q13          1.13      1.38
## q21          1.13      1.38
## q23          1.07      1.17
## q31          1.09      1.26
## q32          1.07      1.17
## Root.P.1.    1.17      1.19
## Root.P.2.    1.23      1.26
## Root.P.3.    1.23      1.27

```

```

# Density Plots
densityplot(res[, -2])

```



```
# Parameter summary
summary(res)
```

```
##
## Iterations = 1100:1e+05
## Thinning interval = 100
## Number of chains = 3
## Sample size per chain = 990
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean      SD Naive SE Time-series SE
## Lh          -42.2394 1.976e+00 3.626e-02    4.033e-01
## Harmonic.Mean -44.1566 2.453e+00 4.500e-02    2.404e-01
## q12          157.7984 9.536e+01 1.750e+00    1.889e+01
## q13          480.1138 2.506e+02 4.599e+00    8.399e+01
## q21          480.1138 2.506e+02 4.599e+00    8.399e+01
## q23           41.6890 4.431e+01 8.130e-01    5.251e+00
## q31          157.7984 9.536e+01 1.750e+00    1.889e+01
## q32           41.6890 4.431e+01 8.130e-01    5.251e+00
## Root.P.1.      0.3333 6.490e-06 1.191e-07    1.652e-07
```

```

## Root.P.2.      0.3333 6.054e-05 1.111e-06      1.571e-06
## Root.P.3.      0.3333 5.441e-05 9.984e-07      1.415e-06
##
## 2. Quantiles for each variable:
##
##              2.5%      25%      50%      75%      97.5%
## Lh            -47.8739 -42.9565 -42.0559 -40.9590 -39.2463
## Harmonic.Mean -48.5052 -46.9362 -43.0291 -42.2816 -40.7952
## q12           20.0634  75.6197 154.8932 223.7319 374.0282
## q13           82.6910 254.3228 497.0869 652.5884 959.1860
## q21           82.6910 254.3228 497.0869 652.5884 959.1860
## q23            1.4314  10.7617  26.6740  56.2958 176.0072
## q31           20.0634  75.6197 154.8932 223.7319 374.0282
## q32            1.4314  10.7617  26.6740  56.2958 176.0072
## Root.P.1.      0.3333   0.3333   0.3333   0.3333   0.3333
## Root.P.2.      0.3333   0.3333   0.3333   0.3333   0.3333
## Root.P.3.      0.3333   0.3333   0.3333   0.3333   0.3333

```

We can conclude two things from this analysis. First, the posterior distribution is strongly affected by the prior used. This is problematic and it suggest that there may not be enough information in the data to properly estimate the transition rate parameters. This is one of the advantage of the Bayesian approach as you can more easily see when it is the case. Second, the chains have not converged as well with the flat prior. This also likely reflect the little information present in the data. Consequently, you should interpret these results with much caution (if at all!).

## Corelated evolution between binary traits in BayesTraits

A common application of phylogenetic methods is to study the correlation of characters (and their evolution). A very popular model for binary traits is that of Pagel (1994). The idea is to test if two traits evolved in a correlated fashion or independently.

The test is run using the Discrete function of BayesTraits. The example below will focus on a Bayesian approach, but this can also be done with ML. The idea is to evaluate two models: one in which the traits evolve independently and another one where the traits evolved in a correlative way.

### Independent model

The simpler model is the independent one. In this model, there are four paramters:

Parameter	Trait	Transitions
$\alpha_1$	1	$0 \rightarrow 1$
$\beta_1$	1	$1 \rightarrow 0$
$\alpha_2$	2	$0 \rightarrow 1$
$\beta_2$	2	$1 \rightarrow 0$

This can be represented in a double transition matrix:

	0,0	0,1	1,0	1,1
0,0	-	$\alpha_2$	$\alpha_1$	0
0,1	$\beta_2$	-	0	$\alpha_1$

	0,0	0,1	1,0	1,1
1,0	$\beta_1$	0	-	$\alpha_2$
1,1	0	$\beta_1$	$\beta_2$	-

Note that the transitions where both characters would have to evolve at the same time are set to zero as this is impossible in an infinitesimal amount of time.

## Dependent model

The dependent model is more complex. It assumes that the rate of change in one character depends on the state of the other character.

Parameter	Dependent on	Trait	Transitions
$q_{1,2}$	Trait 1 = 0	2	$0 \rightarrow 1$
$q_{1,3}$	Trait 2 = 0	1	$0 \rightarrow 1$
$q_{2,1}$	Trait 1 = 0	2	$1 \rightarrow 0$
$q_{2,4}$	Trait 2 = 1	1	$0 \rightarrow 1$
$q_{3,1}$	Trait 2 = 0	1	$1 \rightarrow 0$
$q_{3,4}$	Trait 1 = 1	2	$0 \rightarrow 1$
$q_{4,2}$	Trait 2 = 1	1	$1 \rightarrow 0$
$q_{4,3}$	Trait 1 = 1	2	$1 \rightarrow 0$

As you can see, it has 8 parameters instead of 4. This model results in the following double transition matrix:

	0,0	0,1	1,0	1,1
0,0	-	$q_{1,2}$	$q_{1,3}$	0
0,1	$q_{2,1}$	-	0	$q_{2,4}$
1,0	$q_{3,1}$	0	-	$q_{3,4}$
1,1	0	$q_{4,2}$	$q_{4,3}$	-

Now, BayesTraits can be used to calculate the fit of the two models and then compare them to select the best one.

## Running the analysis

The idea is to run both models separately and compare their fit. In the Bayesian framework, one uses Bayes Factors to compare the models. Let's first fit the model. For this, we will use 500 trees sampled from the posterior distribution of trees. This has the advantage that the analysis will also integrate phylogenetic uncertainty in the model. This is especially important if the support for the groups in the tree are not all very strong. By doing so, the results obtained integrate over all possible tree topology and accounts for phylogenetic uncertainty.

We'll have to import these trees. Such posterior samples of trees can be obtained using Bayesian phylogenetic methods (BEAST, MrBayes).

```
pdtrees <- read.nexus("./data/pd_500.trees")
species.to.exclude <- pdtrees[[1]]$tip.label[!(pdtrees[[1]]$tip.label %in%
rownames(seedplantsdata))]
```

```

pdtrees<-lapply(pdtrees,drop.tip,tip=species.to.exclude)
class(pdtrees)<-"multiPhylo"
attr(pdtrees,"TipLabel") <- pdtrees[[1]]$tip.label
rm(species.to.exclude)
# Need two binary variables
# Start by converting the height variable in a binary variable (0/1)
height2 <- as.numeric(seedplantsdata$height)
height2[height2==2] <- 3
height2[height2==3] <- 0
thedata<-data.frame(code=rownames(seedplantsdata),
                    height=height2,ShadeTolerance=as.numeric(seedplantsdata$ShadeTol)-1)

```

Now, we can perform the BayesTraits analyses using this distribution of trees. For this, we will use the DiscreteMCMC function. The settings are as for above, with a flat prior. The independent and dependent models can be set using the parameter `dependent=FALSE` or `dependent=TRUE`.

```

# Fit independent model
ind.res1 <- DiscreteMCMC(pdtrees, thedata, dependent = FALSE, it = 100000, bi = 1000,
                        sa = 100, rd = 2, pa = "uniform 0 100", silent = TRUE)
ind.res2 <- DiscreteMCMC(pdtrees, thedata, dependent = FALSE, it = 100000, bi = 1000,
                        sa = 100, rd = 2, pa = "uniform 0 100", silent = TRUE)
# Fit dependent model
dep.res1 <- DiscreteMCMC(pdtrees, thedata, dependent = TRUE, it = 100000, bi = 1000,
                        sa = 100, rd = 2, pa = "uniform 0 100", silent = TRUE)
dep.res2 <- DiscreteMCMC(pdtrees, thedata, dependent = TRUE, it = 100000, bi = 1000,
                        sa = 100, rd = 2, pa = "uniform 0 100", silent = TRUE)

```

Now, let's read the results with the coda package.

```

require(coda)
require(lattice)
# Read the BayesTrait results of the independent model in coda format
ind1 <- mcmc(ind.res1$Results[,c(-1,-4)],start=min(ind.res1$Results$Iteration),
            end=max(ind.res1$Results$Iteration),thin=100)
ind2 <- mcmc(ind.res2$Results[,c(-1,-4)],start=min(ind.res2$Results$Iteration),
            end=max(ind.res2$Results$Iteration),thin=100)
# Combine the three chains
ind <- mcmc.list(ind1,ind2)

# Read the BayesTrait results of the dependent model in coda format
dep1 <- mcmc(dep.res1$Results[,c(-1,-4)],start=min(dep.res1$Results$Iteration),
            end=max(dep.res1$Results$Iteration),thin=100)
dep2 <- mcmc(dep.res2$Results[,c(-1,-4)],start=min(dep.res2$Results$Iteration),
            end=max(dep.res2$Results$Iteration),thin=100)
# Combine the three chains
dep <- mcmc.list(dep1,dep2)

```

Now, we can look for convergence of the two sets of analyses

```

# Get effective sizes (should be > 200)
effectiveSize(ind)

```

```
##           Lh Harmonic.Mean           q12           q13           q21
##      216.77551      38.89503      330.40364      394.00370      335.11918
##           q24           q31           q34           q42           q43
##      394.00370      375.17067      330.40364      375.17067      335.11918
## Root...P.0.0. Root...P.0.1. Root...P.1.0. Root...P.1.1.
##      1450.89378      1033.66372      1035.61964      1446.85109
```

```
effectiveSize(dep)
```

```
##           Lh Harmonic.Mean           q12           q13           q21
##      234.76171      14.53585      237.15719      228.40602      298.53284
##           q24           q31           q34           q42           q43
##      225.58948      217.46602      166.88288      239.16590      220.55682
## Root...P.0.0. Root...P.0.1. Root...P.1.0. Root...P.1.1.
##      817.93370      827.04936      851.72739      663.32531
```

```
# Gelman and Rubin's convergence diagnostic
gelman.diag(ind,autoburnin=FALSE,multivariate=FALSE)
```

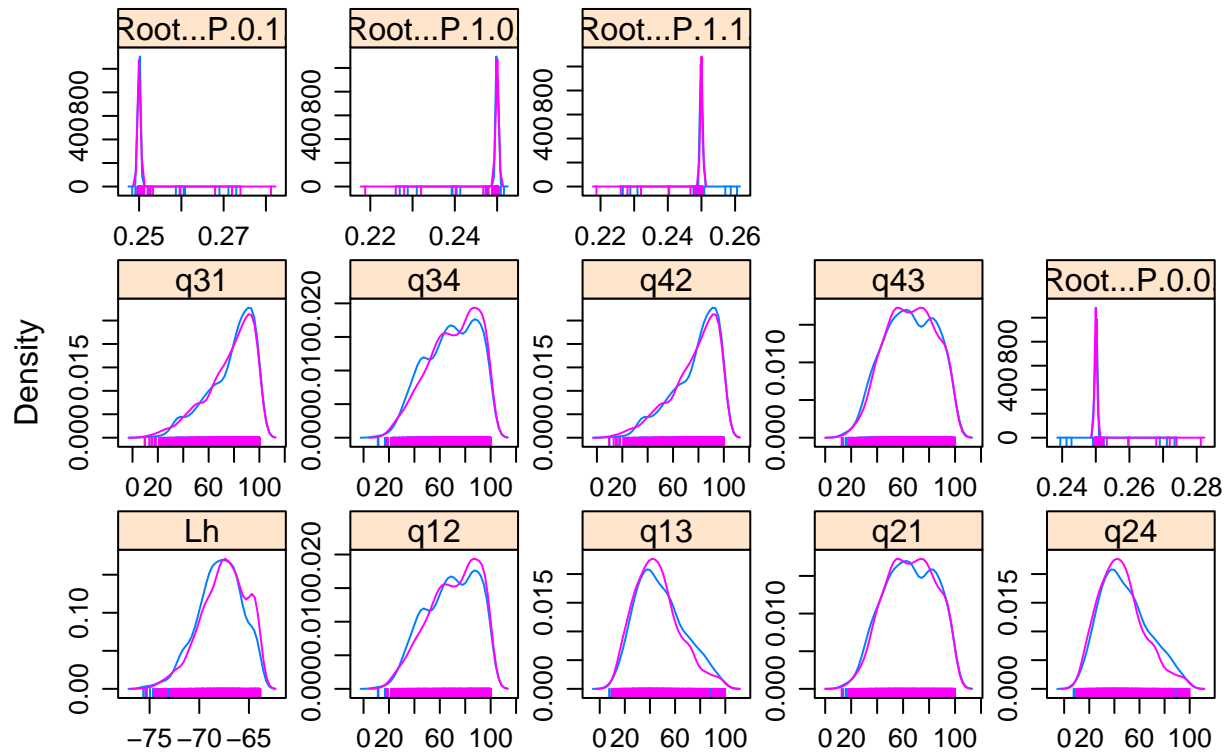
```
## Potential scale reduction factors:
##
##           Point est. Upper C.I.
## Lh           1.01      1.06
## Harmonic.Mean 1.20      1.63
## q12           1.00      1.02
## q13           1.01      1.04
## q21           1.00      1.00
## q24           1.01      1.04
## q31           1.00      1.01
## q34           1.00      1.02
## q42           1.00      1.01
## q43           1.00      1.00
## Root...P.0.0. 1.02      1.02
## Root...P.0.1. 1.01      1.01
## Root...P.1.0. 1.01      1.01
## Root...P.1.1. 1.02      1.02
```

```
gelman.diag(dep,autoburnin=FALSE,multivariate=FALSE)
```

```
## Potential scale reduction factors:
##
##           Point est. Upper C.I.
## Lh           1.03      1.13
## Harmonic.Mean 2.50      5.38
## q12           1.02      1.09
## q13           1.01      1.03
## q21           1.01      1.06
## q24           1.00      1.00
## q31           1.00      1.01
## q34           1.01      1.04
## q42           1.00      1.00
## q43           1.05      1.20
```

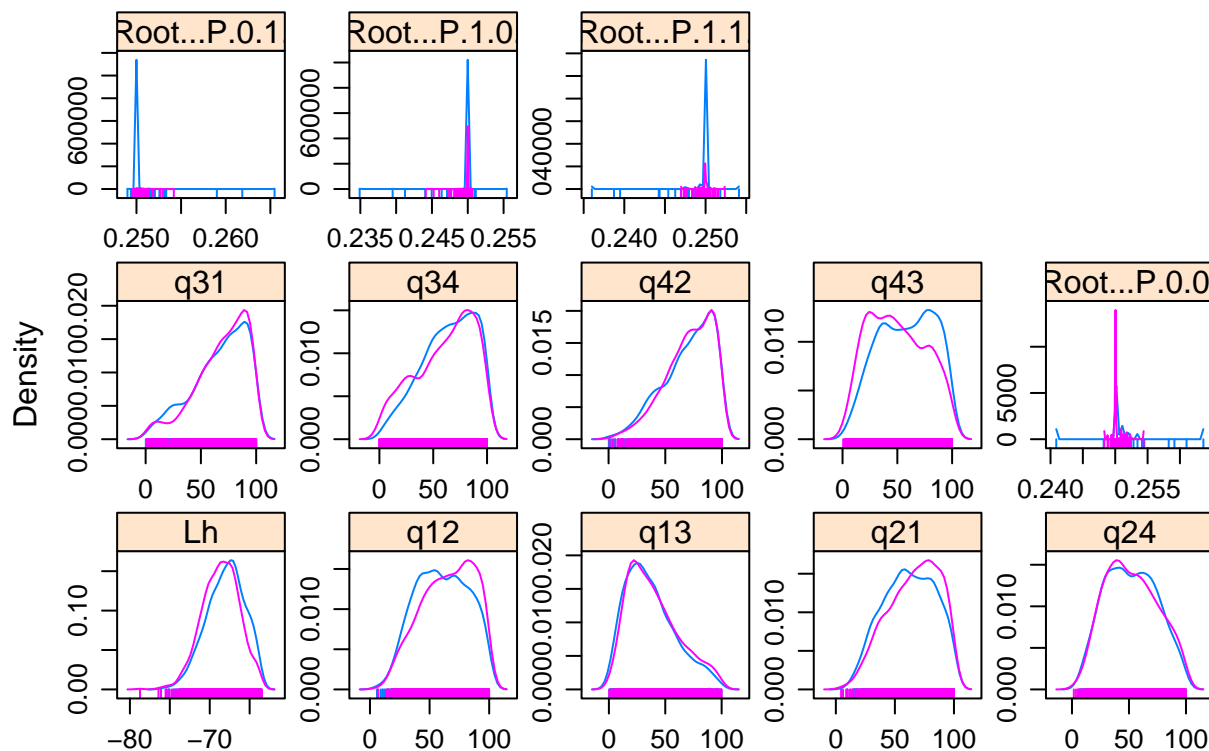
```
## Root...P.0.0.      1.23      1.31
## Root...P.0.1.      1.22      1.24
## Root...P.1.0.      1.16      1.16
## Root...P.1.1.      1.23      1.34
```

```
# Density Plots
densityplot(ind[, -2])
```



```
densityplot(dep[, -2])
```





```
# Parameter summary
summary(ind)
```

```
##
## Iterations = 1100:1e+05
## Thinning interval = 100
## Number of chains = 2
## Sample size per chain = 990
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##              Mean          SD Naive SE Time-series SE
## Lh             -67.7520   2.261793 5.083e-02   1.594e-01
## Harmonic.Mean -70.6596   0.624675 1.404e-02   9.474e-02
## q12             69.8698  19.737675 4.436e-01   1.098e+00
## q13             47.0042  18.695147 4.201e-01   9.479e-01
## q21             65.2890  19.105556 4.294e-01   1.054e+00
## q24             47.0042  18.695147 4.201e-01   9.479e-01
## q31             77.1914  18.424610 4.141e-01   9.516e-01
## q34             69.8698  19.737675 4.436e-01   1.098e+00
## q42             77.1914  18.424610 4.141e-01   9.516e-01
## q43             65.2890  19.105556 4.294e-01   1.054e+00
## Root...P.0.0.   0.2501   0.001450 3.258e-05   4.614e-05
## Root...P.0.1.   0.2501   0.001457 3.274e-05   4.891e-05
## Root...P.1.0.   0.2499   0.001457 3.274e-05   4.888e-05
## Root...P.1.1.   0.2499   0.001449 3.256e-05   4.615e-05
##
## 2. Quantiles for each variable:
```

```
##
##          2.5%    25%    50%    75%    97.5%
## Lh        -72.5325 -69.20 -67.60 -66.11 -64.1023
## Harmonic.Mean -71.1332 -70.93 -70.78 -70.55 -69.8445
## q12        30.3780  55.44  71.65  86.33  98.7847
## q13        16.8409  33.16  44.67  58.15  89.5381
## q21        29.3801  50.54  65.71  80.84  97.1103
## q24        16.8409  33.16  44.67  58.15  89.5381
## q31        34.1042  66.09  82.17  91.98  99.3152
## q34        30.3780  55.44  71.65  86.33  98.7847
## q42        34.1042  66.09  82.17  91.98  99.3152
## q43        29.3801  50.54  65.71  80.84  97.1103
## Root...P.0.0.  0.2500  0.25  0.25  0.25  0.2501
## Root...P.0.1.  0.2500  0.25  0.25  0.25  0.2502
## Root...P.1.0.  0.2498  0.25  0.25  0.25  0.2500
## Root...P.1.1.  0.2499  0.25  0.25  0.25  0.2500
```

```
summary(dep)
```

```
##
## Iterations = 1100:1e+05
## Thinning interval = 100
## Number of chains = 2
## Sample size per chain = 990
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##          Mean          SD Naive SE Time-series SE
## Lh        -68.3337  2.355e+00  5.293e-02    1.526e-01
## Harmonic.Mean -71.3288  7.317e-01  1.644e-02    1.570e-01
## q12        62.8700  2.179e+01  4.897e-01    1.409e+00
## q13        38.3978  2.241e+01  5.036e-01    1.483e+00
## q21        63.8917  2.147e+01  4.825e-01    1.239e+00
## q24        52.1934  2.245e+01  5.044e-01    1.495e+00
## q31        67.6367  2.395e+01  5.382e-01    1.628e+00
## q34        62.2672  2.575e+01  5.786e-01    2.020e+00
## q42        70.3658  2.142e+01  4.813e-01    1.418e+00
## q43        54.0018  2.570e+01  5.777e-01    1.709e+00
## Root...P.0.0.  0.2501  6.003e-04  1.349e-05    2.210e-05
## Root...P.0.1.  0.2501  5.406e-04  1.215e-05    1.887e-05
## Root...P.1.0.  0.2499  5.850e-04  1.315e-05    2.082e-05
## Root...P.1.1.  0.2500  5.761e-04  1.295e-05    2.181e-05
##
## 2. Quantiles for each variable:
##
##          2.5%    25%    50%    75%    97.5%
## Lh        -73.2256 -69.95 -68.23 -66.67 -64.1233
## Harmonic.Mean -72.4960 -71.67 -71.25 -71.00 -68.8548
## q12        22.0143  45.59  63.74  81.33  97.5985
## q13         5.6621  20.77  34.23  52.04  89.7767
## q21        22.3467  47.63  65.34  81.28  98.1083
## q24        13.8632  34.20  51.33  69.45  94.9974
## q31        10.3138  52.90  72.13  87.54  98.4383
```

```
## q34          7.2938  44.61  65.98  83.60  98.2063
## q42         22.0073  56.48  74.34  88.36  98.7528
## q43          9.6942  32.81  53.19  76.63  97.4442
## Root...P.0.0.  0.2499   0.25   0.25   0.25   0.2504
## Root...P.0.1.  0.2500   0.25   0.25   0.25   0.2504
## Root...P.1.0.  0.2494   0.25   0.25   0.25   0.2500
## Root...P.1.1.  0.2497   0.25   0.25   0.25   0.2501
```

You can see that the chains have converged well for both models.

## Bayes Factors

To calculate whether there is support for the more complex correlated model, we will use Bayes Factors, which is common for Bayesian analyses. The Bayes Factor (BF) can be calculated the following way:

$$2\ln BF = 2(\ln L_{\text{complex model}} - \ln L_{\text{simpler model}})$$

To calculate the BF, it is common to use the harmonic mean of the likelihood of each run (but see below). For this, you only use the last value from the complete run.

```
# Harmonic mean of the independent model
(ind_harm<-ind.res1$Results$Harmonic.Mean[length(ind.res1$Results$Harmonic.Mean)])
```

```
## [1] -70.87435
```

```
# Harmonic mean of the dependent model
(dep_harm<-dep.res1$Results$Harmonic.Mean[length(ind.res1$Results$Harmonic.Mean)])
```

```
## [1] -71.02755
```

```
# Bayes Factor
BF = 2*(dep_harm-ind_harm)
BF
```

```
## [1] -0.306388
```

Following Kass and Raftery (1995), BayesFactors can be interpreted the following way:

2 ln BF	Interpretation
0 to 2	Not worth more than a mention
2 to 6	Positive evidence
6 to 10	Strong evidence
> 10	Very strong evidence

Consequently, you can see that with the present case, there is not support for the more complex model.

The harmonic mean is not a very good estimator of the likelihood of the model and many suggest it should not be used. BayesTraits has a stepping stone function to better estimate the likelihood

and if you plan publishing using Bayes Factors, this is what you should use. Unfortunately, there is no wrapper yet to run this function from R.

## Assignment

Please reconstruct the ancestral states of one character on your tree using the method of your choice. Present the analyses and the results in a R Markdown document.

## References

Kass R.E., A.E. Raftery. 1995. Bayes factor. *Journal of the American Statistical Association* 90:773–795.

Pagel M. 1994. Detecting Correlated Evolution on Phylogenies: A General Method for the Comparative Analysis of Discrete Characters. *Proceedings of the Royal Society B* 255:37–45.

Pagel M., A. Meade, D. Barker. 2004. Bayesian estimation of ancestral character states on phylogenies. *Systematic Biology*. 53:673–684.

Paquette A., S. Joly, C. Messier. 2015. Explaining forest productivity using tree functional traits and phylogenetic information: two sides of the same coin over evolutionary scale? *Ecology and Evolution* 5:1774–1783.